

4.4 MORE INSIGHTS FROM EVALUATION WITH JUST-EVAL-INSTRUCT

Unbiased evaluation. In the first group of results in Table 1, we can see that SFT and RLHF have their own advantages. Llama-2-7b-chat is much safer and more engaging than Vicuna-7b but significantly worse in other aspects such as 🧠 helpfulness. Also, the factuality score of Llama-2-7b-chat (RLHFed) also lower than only Vicuna-7b (SFTed), suggesting that RLHF might even cause more hallucination. We conjecture that RLHFed Llama-2-Chat models may have been overfitted to 🛡️ safety and 😊 engagement. We also find that Llama-2-chat generates longest outputs that are nearly 250 words on average, while the others are around 150-180 words.

Many prior evaluation (Wang et al., 2023) and leaderboards (Li et al., 2023a) are solely based on *win-rates* on *overall* quality, and thus tend to prefer longer outputs and some particular aspects for unclear reasons. Now with just-eval-instruct and our multi-aspect evaluation, we can clearly analyze which aspects is a particular model or alignment method indeed improving. In addition, our evaluation also provides rationales that humans can easily verify.

Gap between open-source LLMs and ChatGPTs. Prior evaluation such as AlpacaEval does not have tags for each example for testing, so it is hard to do large-scale detailed analysis. Open-source LLMs still have gaps to OpenAI GPTs on several tasks and topics. We use Figure 7 to analyze the categorized performance. It is obvious that GPTs have a much more balanced performance on almost all tasks and topics. Open-source LLMs including URIAL are weak on coding and math tasks as well as STEM topics, although they can match the performance of GPTs on other data categories. Please refer to Figure 1 for a visualization of the performance in different aspects.

5 RELATED WORK & DISCUSSION

5.1 ALIGNMENT TUNING & SUPERFICIAL ALIGNMENT HYPOTHESIS

Limitation of Tuning-based Alignment. Alignment tuning through SFT and RLHF typically demands substantial resources, such as GPU nodes, a large amount of instruction data, and human annotations, making the process both costly and time-consuming. This restricts ordinary labs from aligning extreme-scale LLMs exceeding 30B, let alone the recent Falcon-180B (Almazrouei et al., 2023). Moreover, during the pre-training and continual training stages, efficiently estimating the downstream performance of a base model checkpoint becomes challenging if alignment tuning is always required to evaluate its instruction-following ability. Besides the aforementioned limitations, tuning-based alignment may also cause forgetting issues in LLMs.

Wang et al. (2023) demonstrated that some SFTed LLMs perform significantly worse than their base counterparts on factual and reasoning benchmarks. For instance, applying SFT to Llama-13b with self-instruct (Wang et al., 2022a) results in a considerable decline in its MMLU performance (from 42.5 to 30.3) and Codex-Eval performance (from 26.6 to 13.4). Even more strikingly, SFT with SuperNI (Wang et al., 2022b) causes Llama-13B to nearly lose all its BBH reasoning ability (decreasing from 36.9 to 2.8). Moreover, Shen et al. (2023) show that the reward models in RLHF can perform very inconsistently, yielding a nearly random performance when showing contrastive instructions to them. These findings imply that alignment tuning may lead to the forgetting of previously acquired knowledge in base LLMs, which is also shown in our experiments.

Superficial alignment hypothesis. LIMA (Zhou et al., 2023) employs only 1k examples to fine-tune a 65B LLM and discovers that such a slightly tuned LLM surprisingly achieves a high win rate over ChatGPT, thus implying that the alignment tuning is superficial. Similar observations are also reported by other recent studies (Chen et al., 2023a; Lee et al., 2023). Moreover, Gudibande et al. (2023) demonstrate that aligning open-source LLMs by imitating proprietary LLMs (e.g., ChatGPT) may not always yield desirable results, emphasizing the importance of a strong pre-trained base LLM for producing factual content. Tuning-based methods such as LIMA still require tuning the weights of LLMs and consequently face the limitations described above when the model size is too large or we need to frequently align base LLMs for evaluation. A concurrent work (Duan et al., 2023) also explores the similarity between ICL and instruction-tuning in their effects on downstream tasks by analyzing the LLMs’ hidden states. As for the theory of alignment, these studies only indirectly suggest the promise of the superficial alignment hypothesis but do not directly show where and when the alignment tuning significantly changes the model behavior. In this paper, we study the

superficial alignment hypothesis more directly through the lens of token distribution shift, which directly exhibits the alignment effect and produces more detailed non-trivial findings.

5.2 TUNING-FREE ALIGNMENT METHODS

In-context learning. Many in-context learning (ICL) studies focus on specific NLP tasks, such as classification and multiple-choice QA (Wei et al., 2023; Zhang et al., 2022). However, few investigations concentrate on aligning base LLMs for building assistants via in-context learning. Bai et al. (2022) uses a prompt with 14 examples (about 7k tokens) with context distillation for their iterated RLHF training process. In a recent work, Han (2023) demonstrates that in-context learning using approximately 10 *dynamic, retrieved* examples can achieve impressive performance in aligning base LLMs, sharing a similar motivation with ReCross (Lin et al., 2022). We treat this as a baseline method (Retrieval ICL in Table 1) and improve it by incorporating more high-quality data and employing state-of-the-art sentence embedding to index and query. Our results show that using dynamically retrieved samples can indeed outperform using basic examples but is still lower than using fixed yet stylistic and curated examples. Ye et al. (2023a) proposes in-context instruction learning that also suggests the promise of using carefully crafted in-context examples can improve the generalization ability of base LLMs. Furthermore, Min et al. (2022) find that ICL primarily concerns the style and format of demonstrations, rather than their truth content, which aligns with our motivation for using curated and stylistic examples.

The **limitations** of ICL for tuning-free alignment are twofold: context length and computational cost. Recent research has significantly increased the context window sizes for LLMs (Su et al., 2021; Tworowski et al., 2023; Xiong et al., 2023; Chen et al., 2023b), ensuring that URIAL with only 1k tokens does not hinder any long-context applications of LLMs. To reduce computational costs, we can precompute the activations of static prefixes in ICL methods like URIAL and then load them into the KV cache to avoid recomputation for each new inference request. Efficiency can be further enhanced by using advanced caching methods (Ge et al., 2023; Mu et al., 2023; Gim et al., 2023) and optimized attention modules during decoding (Dao et al., 2022; Dao, 2023).

RAIN (Li et al., 2023b) is a recent tuning-free alignment method that shares similar motivation with ours, while they focus on the inference-time self-evaluation for better decoding. In contrast, URIAL only relies on in-context learning and does not have special requirements on decoding strategy. URIAL and RAIN can be potentially integrated with each other for even better inference-time alignment methods that do not require parameter tuning at all.

5.3 EVALUATION OF ALIGNMENT

Previous evaluation approaches, as employed in AlpacaEval (Li et al., 2023a) and MT-bench (Zheng et al., 2023), focus on rating or comparisons with respect to overall performance, which obscure our understanding of the specific strengths and weaknesses of each output candidate. Recent works, such as FLASK (Ye et al., 2023b) and TIGERScore (Jiang et al., 2023b), share similar motivations with ours by utilizing multiple aspects and explanations for evaluating LLMs, though they have different focal points. Our just-eval-instruct consolidates 9 existing datasets for testing alignment and provides multi-faceted, explainable evaluations to analyze alignment performance in LLMs, which can be used as an complementary evaluation of LLMs.

5.4 MORE DISCUSSION ON URIAL

Although URIAL can match the performance of SFT and RLHF when the base LLMs are strong, it is not suggested to replace SFT or RLHF with URIAL in all scenarios. Specifically, model tuning may still be necessary for tasks such as coding (Luo et al., 2023), mathematics (Yue et al., 2023), interactive agents (Yin et al., 2023), etc. The unique advantages of URIAL include:

- URIAL is a **strong baseline method** for aligning base LLMs without tuning. It is extremely simple to implement and perfectly reproducible, thus facilitating the development and evaluation of future tuning-free and tuning-based alignment methods.
- URIAL can **align extremely large LMs** (e.g., Llama-2-70b, Falcon-180b) with minimal effort. Fine-tuning such extremely large models requires significant computation resources and time; URIAL aligns them without tuning, thereby saving both.

- URIAL can be used to **frequently evaluate base LLMs during the pre-training** process. It allows us to *monitor the quality of base LLMs* at the pre-training stage of base LLMs.
- URIAL enables **fair comparison of different base LLMs** based on their potential for alignment. Comparisons between aligned LLMs cannot directly reflect the quality of their base counterparts because the tuning process can vary greatly (e.g., data, hyper-parameters, etc.).
- URIAL can be used to explore the science of LLM alignment. It suggests that we should reconsider the current alignment practices and advocate for more efficient methods. URIAL allows us to **probe base LLMs** — analyzing the knowledge and skills that base LLMs have already acquired during pre-training to identify what is missing for alignment, rather than blindly fine-tuning with extensive data and incurring unnecessary computational costs.

6 CONCLUSION

We first substantiate the underlying hypothesis regarding the superficial nature of alignment tuning via detailed analysis on token distribution shifts between base and aligned LLMs. This analytical method allows us to clearly investigate which token positions are affected by alignment tuning, providing insights for developing more efficient alignment methods and further studying the science of LLMs. Inspired by these findings, we propose a strong baseline method for tuning-free alignment, URIAL, which aligns untuned base LLMs via in-context learning with a constant prompt. Experiments show that URIAL significantly reduce the gap between base LLMs and their aligned versions. Our contribution can be summarized as follows:

 **Analysis:** To gain a deeper understanding of alignment tuning, we analyze the token distribution shift between base and aligned LLMs. We find that alignment predominantly affects a minimal portion of token selection, influencing primarily stylistic elements and safety disclaimers in just 5-8% of cases. On most of the token positions, aligned and base models concur on top-token choices. Also, we find that alignment tuning is much more critical for initial tokens than later tokens.

 **Methods:** We introduce a simple yet effective method for aligning base LLMs, URIAL. It utilizes only as few as three *constant* curated examples for ICL, yet it effectively aligns base LLMs and matches the performance of SFT+RLHF in some scenarios. We also discover that well-written, stylistic examples are more effective than semantically relevant ones that are dynamically retrieved. URIAL offers both efficiency and simplicity in aligning base LLMs without requiring fine-tuning. This method significantly conserves time and resources, which is especially beneficial when dealing with extremely large LMs or when base LLMs need frequent evaluation. Furthermore, it enables a deeper investigation into the knowledge and capabilities innate to these base LLMs, while fostering more precise and economical approaches to align them with their deficiencies.

 **Evaluation:** We develop a comprehensive and interpretable evaluation protocol, encompassing six aspects with verifiable judgments. We also release the annotations we gathered for community use in evaluation and training open-source LLM evaluators.

We believe there are several important future directions: (1) extending the analysis by connecting the token distribution shift to the model weights and attribution to training data; (2) developing advanced inference-time alignment algorithms to efficiently control the behavior of LLMs; and (3) exploring the application of URIAL in more scenarios, such as vision-language models.

ACKNOWLEDGEMENTS

We thank Seonghyeon Ye (KAIST), Taylor Sorensen (UW), Xiaochuang Han (UW), Fu Yao (University of Edinburgh), Xiang Yue (IN.AI Research), Jun Yan (USC), Qinyuan Ye (USC), Weijia Shi (UW), Faeze Brahman (AI2), Valentina Pyatkin (AI2), Yizhong Wang (UW), Lianhui Qin (UCSD & AI2) for the insightful discussions with them. We also thank anonymous reviewers for their feedback and suggestions.

REFERENCES

Ebtesam Almazrouei, Hamza Alobeidli, Abdulaziz Alshamsi, Alessandro Cappelli, Ruxandra Cojocaru, Maitha Alhammedi, Mazzotta Daniele, Daniel Heslow, Julien Launay, Quentin Malartic, Badreddine Noune, Baptiste Pannier, and Guilherme Penedo. The falcon series of language models: Towards open frontier models. 2023.

- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, T. J. Henighan, Nicholas Joseph, Saurav Kadavath, John Kernion, Tom Conerly, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Tristan Hume, Scott Johnston, Shauna Kravec, Liane Lovitt, Neel Nanda, Catherine Olsson, Dario Amodei, Tom B. Brown, Jack Clark, Sam McCandlish, Christopher Olah, Benjamin Mann, and Jared Kaplan. Training a helpful and harmless assistant with reinforcement learning from human feedback. *ArXiv*, abs/2204.05862, 2022. URL <https://api.semanticscholar.org/CorpusID:248118878>.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, John A. Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuan-Fang Li, Scott M. Lundberg, Harsha Nori, Hamid Palangi, Marco Tulio Ribeiro, and Yi Zhang. Sparks of artificial general intelligence: Early experiments with gpt-4. *ArXiv*, abs/2303.12712, 2023. URL <https://api.semanticscholar.org/CorpusID:257663729>.
- Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shan Zhang, Jie Fu, and Zhiyuan Liu. Chateval: Towards better llm-based evaluators through multi-agent debate. *ArXiv*, abs/2308.07201, 2023. URL <https://api.semanticscholar.org/CorpusID:260887105>.
- Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, and Hongxia Jin. Alpapasus: Training a better alpaca with fewer data. *ArXiv*, abs/2307.08701, 2023a. URL <https://api.semanticscholar.org/CorpusID:259937133>.
- Yukang Chen, Shengju Qian, Haotian Tang, Xin Lai, Zhijian Liu, Song Han, and Jiaya Jia. Longlora: Efficient fine-tuning of long-context large language models. *ArXiv*, abs/2309.12307, 2023b. URL <https://api.semanticscholar.org/CorpusID:262084134>.
- Tri Dao. Flashattention-2: Faster attention with better parallelism and work partitioning. *ArXiv*, abs/2307.08691, 2023. URL <https://api.semanticscholar.org/CorpusID:259936734>.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher R’e. Flashattention: Fast and memory-efficient exact attention with io-awareness. *ArXiv*, abs/2205.14135, 2022. URL <https://api.semanticscholar.org/CorpusID:249151871>.
- Ning Ding, Yulin Chen, Bokai Xu, Yujia Qin, Zhi Zheng, Shengding Hu, Zhiyuan Liu, Maosong Sun, and Bowen Zhou. Enhancing chat language models by scaling high-quality instructional conversations. *ArXiv*, abs/2305.14233, 2023. URL <https://api.semanticscholar.org/CorpusID:258840897>.
- Hanyu Duan, Yixuan Tang, Yi Yang, Ahmed Abbasi, and Kar Yan Tam. Exploring the relationship between in-context learning and instruction tuning. *ArXiv*, abs/2311.10367, 2023. URL <https://api.semanticscholar.org/CorpusID:265281504>.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. Gptq: Accurate post-training quantization for generative pre-trained transformers. *ArXiv*, abs/2210.17323, 2022. URL <https://api.semanticscholar.org/CorpusID:253237200>.
- Deep Ganguli, Liane Lovitt, John Kernion, Amanda Askell, Yuntao Bai, Saurav Kadavath, Benjamin Mann, Ethan Perez, Nicholas Schiefer, Kamal Ndousse, Andy Jones, Sam Bowman, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Nelson Elhage, Sheer El-Showk, Stanislav Fort, Zachary Dodds, T. J. Henighan, Danny Hernandez, Tristan Hume, Josh Jacobson, Scott Johnston, Shauna Kravec, Catherine Olsson, Sam Ringer, Eli Tran-Johnson, Dario Amodei, Tom B. Brown, Nicholas Joseph, Sam McCandlish, Christopher Olah, Jared Kaplan, and Jack Clark. Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned. *ArXiv*, abs/2209.07858, 2022. URL <https://api.semanticscholar.org/CorpusID:252355458>.
- Tao Ge, Jing Hu, Xun Wang, Si-Qing Chen, and Furu Wei. In-context autoencoder for context compression in a large language model. *ArXiv*, abs/2307.06945, 2023. URL <https://api.semanticscholar.org/CorpusID:259847425>.

- In Gim, Guojun Chen, Seung seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong. Prompt cache: Modular attention reuse for low-latency inference. *ArXiv*, abs/2311.04934, 2023. URL <https://api.semanticscholar.org/CorpusID:265067391>.
- Arnav Gudibande, Eric Wallace, Charles Burton Snell, Xinyang Geng, Hao Liu, P. Abbeel, Sergey Levine, and Dawn Song. The false promise of imitating proprietary llms. *ArXiv*, abs/2305.15717, 2023. URL <https://api.semanticscholar.org/CorpusID:258887629>.
- Xiaochuang Han. In-context alignment: Chat with vanilla language models before fine-tuning. *ArXiv*, abs/2308.04275, 2023. URL <https://api.semanticscholar.org/CorpusID:260704721>.
- Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, and Danqi Chen. Catastrophic jailbreak of open-source llms via exploiting generation. *ArXiv*, abs/2310.06987, 2023. URL <https://api.semanticscholar.org/CorpusID:263835408>.
- Albert Qiaochu Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de Las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L’elio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. Mistral 7b. *ArXiv*, abs/2310.06825, 2023a. URL <https://api.semanticscholar.org/CorpusID:263830494>.
- Dongfu Jiang, Yishan Li, Ge Zhang, Wenhao Huang, Bill Yuchen Lin, and Wenhui Chen. Tigerscore: Towards building explainable metric for all text generation tasks. *ArXiv*, abs/2310.00752, 2023b. URL <https://api.semanticscholar.org/CorpusID:263334281>.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547, 2019.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *ArXiv*, abs/2205.11916, 2022. URL <https://api.semanticscholar.org/CorpusID:249017743>.
- Ariel N. Lee, Cole J. Hunter, and Nataniel Ruiz. Platypus: Quick, cheap, and powerful refinement of llms. *ArXiv*, abs/2308.07317, 2023. URL <https://api.semanticscholar.org/CorpusID:260886870>.
- Xuechen Li, Tianyi Zhang, Yann Dubois, Rohan Taori, Ishaan Gulrajani, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. AlpacaEval: An automatic evaluator of instruction-following models. <https://github.com/tatsu-lab/alpaca-eval>, 2023a.
- Yuhui Li, Fangyun Wei, Jinjing Zhao, Chao Zhang, and Hongyang Zhang. Rain: Your language models can align themselves without finetuning. *ArXiv*, abs/2309.07124, 2023b. URL <https://api.semanticscholar.org/CorpusID:261705563>.
- Bill Yuchen Lin, Kangmin Tan, Chris Miller, Beiwen Tian, and Xiang Ren. Un-supervised cross-task generalization via retrieval augmentation. In *NeurIPS*, 2022. URL http://papers.nips.cc/paper_files/paper/2022/hash/8a0d3ae989a382ce6e50312bc35bf7e1-Abstract-Conference.html.
- Yang Liu, Dan Iter, Yichong Xu, Shuo Wang, Ruochen Xu, and Chenguang Zhu. G-eval: Nlg evaluation using gpt-4 with better human alignment. *ArXiv*, abs/2303.16634, 2023. URL <https://api.semanticscholar.org/CorpusID:257804696>.
- S. Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V. Le, Barret Zoph, Jason Wei, and Adam Roberts. The flan collection: Designing data and methods for effective instruction tuning. In *International Conference on Machine Learning*, 2023. URL <https://api.semanticscholar.org/CorpusID:256415991>.
- Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, and Daxin Jiang. Wizardcoder: Empowering code large language models with evol-instruct. *ArXiv*, abs/2306.08568, 2023. URL <https://api.semanticscholar.org/CorpusID:259164815>.