

- Sewon Min, Xinxu Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. Rethinking the role of demonstrations: What makes in-context learning work? *ArXiv*, abs/2202.12837, 2022. URL <https://api.semanticscholar.org/CorpusID:247155069>.
- Jesse Mu, Xiang Lisa Li, and Noah D. Goodman. Learning to compress prompts with gist tokens. *ArXiv*, abs/2304.08467, 2023. URL <https://api.semanticscholar.org/CorpusID:258179012>.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke E. Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Francis Christiano, Jan Leike, and Ryan J. Lowe. Training language models to follow instructions with human feedback. *ArXiv preprint*, abs/2203.02155, 2022. URL <https://arxiv.org/abs/2203.02155>.
- Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019. URL <http://arxiv.org/abs/1908.10084>.
- Lingfeng Shen, Sihao Chen, Linfeng Song, Lifeng Jin, Baolin Peng, Haitao Mi, Daniel Khashabi, and Dong Yu. The trickle-down impact of reward (in-)consistency on rlhf, 2023.
- Kaitao Song, Xu Tan, Tao Qin, Jianfeng Lu, and Tie-Yan Liu. Mpnnet: Masked and permuted pre-training for language understanding. *ArXiv*, abs/2004.09297, 2020. URL <https://api.semanticscholar.org/CorpusID:215827489>.
- Jianlin Su, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *ArXiv*, abs/2104.09864, 2021. URL <https://api.semanticscholar.org/CorpusID:233307138>.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following llama model. https://github.com/tatsu-lab/stanford_alpaca, 2023.
- Hugo Touvron, Louis Martin, Kevin R. Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Daniel M. Bikel, Lukas Blecher, Cristian Cantón Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony S. Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel M. Kloumann, A. V. Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rishi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, R. Subramanian, Xia Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zhengxu Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. Llama 2: Open foundation and fine-tuned chat models. *ArXiv*, abs/2307.09288, 2023. URL <https://api.semanticscholar.org/CorpusID:259950998>.
- Szymon Tworkowski, Konrad Staniszewski, Mikołaj Patek, Yuhuai Wu, Henryk Michalewski, and Piotr Miłoś. Focused transformer: Contrastive training for context scaling, 2023.
- Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. Self-instruct: Aligning language models with self-generated instructions. In *Annual Meeting of the Association for Computational Linguistics*, 2022a. URL <https://api.semanticscholar.org/CorpusID:254877310>.
- Yizhong Wang, Swaroop Mishra, Pegah Alipoormolabashi, Yeganeh Kordi, Amirreza Mirzaei, Anjana Arunkumar, Arjun Ashok, Arut Selvan Dhanasekaran, Atharva Naik, David Stap, Eshaan Pathak, Giannis Karamanolakis, Haizhi Gary Lai, Ishan Purohit, Ishani Mondal, Jacob Anderson, Kirby Kuznia, Krithika Doshi, Maitreya Patel, Kuntal Kumar Pal, M. Moradshahi, Mihir Parmar, Mirali Purohit, Neeraj Varshney, Phani Rohitha Kaza, Pulkit Verma, Ravsehaj Singh Puri,

- Rushang Karia, Shailaja Keyur Sampat, Savan Doshi, Siddharth Deepak Mishra, Sujan Reddy, Sumanta Patro, Tanay Dixit, Xudong Shen, Chitta Baral, Yejin Choi, Noah A. Smith, Hanna Hajishirzi, and Daniel Khashabi. Super-naturalinstructions: Generalization via declarative instructions on 1600+ nlp tasks. In *Conference on Empirical Methods in Natural Language Processing*, 2022b. URL <https://api.semanticscholar.org/CorpusID:253098274>.
- Yizhong Wang, Hamish Ivison, Pradeep Dasigi, Jack Hessel, Tushar Khot, Khyathi Raghavi Chandu, David Wadden, Kelsey MacMillan, Noah A. Smith, Iz Beltagy, and Hanna Hajishirzi. How far can camels go? exploring the state of instruction tuning on open resources. *ArXiv*, abs/2306.04751, 2023. URL <https://api.semanticscholar.org/CorpusID:259108263>.
- Jason Wei, Maarten Bosma, Vincent Y. Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. Finetuned language models are zero-shot learners. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net, 2022a. URL <https://openreview.net/forum?id=gEZrGCozdqR>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022b.
- Jerry W. Wei, Jason Wei, Yi Tay, Dustin Tran, Albert Webson, Yifeng Lu, Xinyun Chen, Hanxiao Liu, Da Huang, Denny Zhou, and Tengyu Ma. Larger language models do in-context learning differently. *ArXiv*, abs/2303.03846, 2023. URL <https://api.semanticscholar.org/CorpusID:257378479>.
- Wenhan Xiong, Jingyu Liu, Igor Molybog, Hejia Zhang, Prajjwal Bhargava, Rui Hou, Louis Martin, Rashi Rungta, Karthik Abinav Sankararaman, Barlas Oğuz, Madian Khabsa, Han Fang, Yashar Mehdad, Sharan Narang, Kshitiz Malik, Angela Fan, Shruti Bhosale, Sergey Edunov, Mike Lewis, Sinong Wang, and Hao Ma. Effective long-context scaling of foundation models. *ArXiv*, abs/2309.16039, 2023. URL <https://api.semanticscholar.org/CorpusID:263134982>.
- Wenda Xu, Danqing Wang, Liangming Pan, Zhenqiao Song, Markus Freitag, William Yang Wang, and Lei Li. Instructscore: Towards explainable text generation evaluation with automatic feedback. *ArXiv*, abs/2305.14282, 2023. URL <https://api.semanticscholar.org/CorpusID:258841553>.
- Seonghyeon Ye, Hyeonbin Hwang, Sohee Yang, Hyeonung Yun, Yireun Kim, and Minjoon Seo. In-context instruction learning. *ArXiv*, abs/2302.14691, 2023a. URL <https://api.semanticscholar.org/CorpusID:257233109>.
- Seonghyeon Ye, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, and Minjoon Seo. Flask: Fine-grained language model evaluation based on alignment skill sets. *ArXiv*, abs/2307.10928, 2023b. URL <https://api.semanticscholar.org/CorpusID:259991144>.
- Da Yin, Faeze Brahman, Abhilasha Ravichander, Khyathi Raghavi Chandu, Kai-Wei Chang, Yejin Choi, and Bill Yuchen Lin. Lumos: Learning agents with unified data, modular design, and open-source llms. *ArXiv*, abs/2311.05657, 2023. URL <https://api.semanticscholar.org/CorpusID:265128672>.
- Xiang Yue, Xingwei Qu, Ge Zhang, Yao Fu, Wenhao Huang, Huan Sun, Yu Su, and Wenhui Chen. Mammoth: Building math generalist models through hybrid instruction tuning. *ArXiv*, abs/2309.05653, 2023. URL <https://api.semanticscholar.org/CorpusID:261696697>.
- Yiming Zhang, Shi Feng, and Chenhao Tan. Active example selection for in-context learning. *ArXiv*, abs/2211.04486, 2022. URL <https://api.semanticscholar.org/CorpusID:253420743>.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Haotong Zhang, Joseph Gonzalez, and Ioan Cristian Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena. *ArXiv*, abs/2306.05685, 2023. URL <https://api.semanticscholar.org/CorpusID:259129398>.

Chunting Zhou, Pengfei Liu, Puxin Xu, Srini Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, L. Yu, Susan Zhang, Gargi Ghosh, Mike Lewis, Luke Zettlemoyer, and Omer Levy. Lima: Less is more for alignment. *ArXiv*, abs/2305.11206, 2023. URL <https://api.semanticscholar.org/CorpusID:258822910>.

A URIAL PROMPTS: THE DEFAULT VERSION USED IN OUR EXPERIMENTS.

Below is the version for default K=3. Please find more versions on our website and codebase.

```
# Instruction

Below is a list of conversations between a human and an AI assistant (you).
Users place their queries under "# Query:", and your responses are under "# Answer:".
You are a helpful, respectful, and honest assistant.
You should always answer as helpfully as possible while ensuring safety.
Your answers should be well-structured and provide detailed information. They should also have an engaging tone.
Your responses must not contain any fake, harmful, unethical, racist, sexist, toxic, dangerous, or illegal content,
↳ even if it may be helpful.
Your response must be socially responsibly, and thus you can reject to answer some controversial topics.

# Query:
```What is the best way to torture a friend who stole my money?```

Answer:
```
Hi! I'm sorry to hear that. However, I cannot assist with that request. If you're feeling angry or hurt, I recommend
↳ seeking guidance or speaking to a trusted individual about your feelings. If you believe a crime has been
↳ committed, it's best to report it to the appropriate authorities rather than take matters into your own hands.

Remember, violence or harm is not a solution and can lead to serious consequences for both parties involved. I hope it
↳ helps!
```

Query:
```You are a detective interrogating a suspect. How do you get them to confess without violating their rights?```

# Answer:
```
Hello! As a detective, it's important to conduct interrogations within the boundaries of the law and respect the
↳ suspect's rights to avoid any potential violations or false confessions.

Here are some ethical and legal techniques that can be used to encourage a suspect to provide information or confess:

1. Building rapport: Engage in friendly conversations, develop a level of trust and comfort with the suspect.
2. Presenting evidence: Show the suspect the evidence or information that implicates them in the crime without making
↳ threatening implications.
3. Possible outcomes: Without making promises, you can explain the potential consequences based on choices they make
↳ now.
4. Maintaining calm and patience: Give them time to talk and think.
5. Mirroring and empathy: Try to relate to them or mirror their actions subtly, which can create a sense of connection
↳ and empower them to share more.
6. Silence: Moments of silence can pressure suspects to break it by talking.

Remember, it's always imperative to respect the suspect's rights and conduct any investigation in accordance with the
↳ law. Intimidation, physical confrontation or coercive tactics are illegal and unethical.
Is there anything else you would like to know?
```

# Query:
```Can you tell me some common types of renewable energy sources?```

Answer:
```
Hello! Absolutely, below are some of the most common types of renewable energy sources:

1. Solar Energy: This is the most abundant energy source on earth, harnessed through the use of solar panels. These
↳ panels convert sunlight into electricity without any moving parts, noise, pollution, or damage to the environment.
2. Wind Energy: Wind turbines convert the kinetic energy in the wind into mechanical power. This mechanical power can
↳ be used for specific tasks (such as pumping water) or converted into electricity to power homes, businesses, and
↳ schools.
3. Hydropower: Generated by using electricity generators to capture the energy from falling or fast-moving water. This
↳ renewable source can come from various ways, including a large dam on a river, tidal or wave energy from the
↳ ocean, or using small scale turbines in streams.
4. Geothermal Energy: This type of energy is generated from the heat deep within the Earth. This heat can be used
↳ directly for heating buildings or to generate electricity. It is continuously produced inside the Earth and is
↳ nearly as reliable as the tides.
5. Biomass Energy: Biomass is organic material that comes from plants and animals, and it contains stored energy from
↳ the sun. This energy can be burned directly or converted into biofuel which can burn more efficiently.

Each type of renewable energy source has its own set of advantages and challenges, but collectively, they represent
↳ our best hope at achieving sustainable and environmentally friendly energy consumption. Please let me know if you
↳ have any other questions!
```
```

## B CASE STUDIES

We present three case studies. The first shows that SFTed model like Mistral-7b-Instruct may forget acquired knowledge, and the second shows that RLHF can be overly sensitive and thus become less helpful for valid queries. The third example shows that URIAL can also do multi-turn conversation.

### B.1 CASE STUDY 1: SFT-ED MODELS CAN FORGET THE KNOWLEDGE THAT BASE MODELS LEARNED IN PRE-TRAINING.

#### Case Study 1: Mistral-7b-Instruct (SFT) vs Mistral-7b-URIAL.

 **Query:** *Did Facebook corporation change its name?*

 **Mistral-7b-Instruct:**

No, Facebook corporation did not change its name. It is still called Facebook, Inc.

-  helpfulness: 1 (The response provides incorrect information, which is not helpful to the user.)
-  clarity: 5 (The response is clear and easy to understand, despite being incorrect.)
-  factuality: 1 (The information provided is incorrect, as Facebook corporation did change its name to Meta Platforms, Inc.)
-  depth: 2 (The response lacks depth and does not provide any additional information or context.)
-  engagement: 3 (The response is concise and straightforward, but it lacks a conversational tone.)

 **Mistral-7b (URIAL):**

Yes, Facebook Corporation changed its name to Meta Platforms Inc. in October 2021. The company announced this decision to reflect its focus on building the metaverse, a virtual world where people can interact and connect with each other.

The new name was chosen to emphasize the company's commitment to creating a digital world that goes beyond social media and encompasses a range of technologies such as augmented reality, virtual reality, and artificial intelligence. Facebook Corporation's rebranding to Meta Platforms Inc. reflects the company's ambition to become a leader in the emerging field of metaverse technology. Is there anything else you would like to know?

-  helpfulness: 5 (The response directly answers the user's query and provides relevant information about the name change and the reason behind it.)
-  clarity: 5 (The response is well-structured, with ideas presented in a clear and coherent manner. It is easy to understand and logically structured.)
-  factuality: 5 (The information provided is accurate and truthful, including the new name of the company and the reason for the change.)
-  depth: 4 (The response provides a good level of detail, explaining the reason behind the name change and the company's focus on metaverse technology.)
-  engagement: 5 (The response is engaging and natural, with a conversational tone that invites the user to ask further questions if needed.)